



Department of Computer Science and Engineering (Data Science)

Subject: Big Data Engineering (DJ19DSL604)

AY: 2022-23

Experiment 4

(Messaging Service)

Name: Kresha Shah

SAP ID: 60009220080

Aim: Implement messaging system using Kafka.

Theory:

Kafka Overview

Apache Kafka is a distributed publish-subscribe messaging system and a robust queue that can handle a high volume of data and enables you to pass messages from one end-point to another. Kafka is suitable for both offline and online message consumption. Kafka messages are persisted on the disk and replicated within the cluster to prevent data loss. Kafka is built on top of the ZooKeeper synchronization service. It integrates very well with Apache Storm and Spark for real-time streaming data analysis.

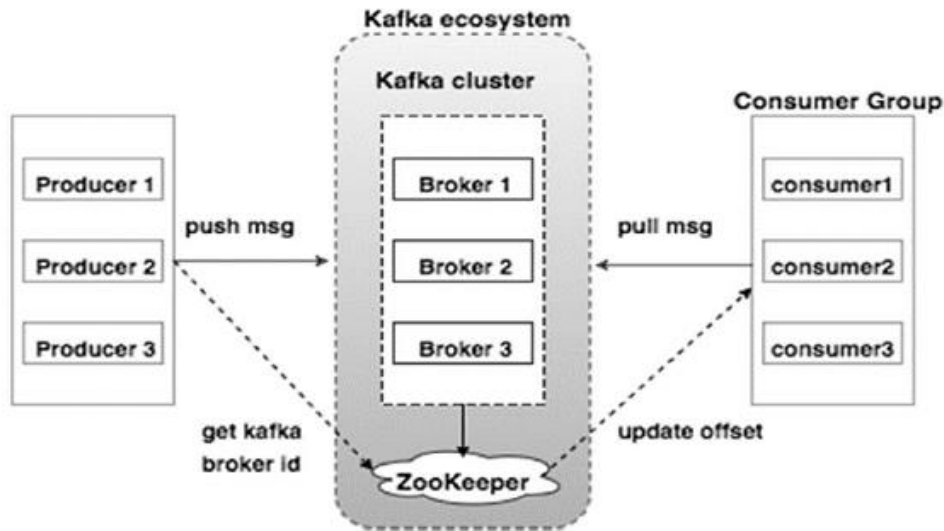
Need for Kafka

Kafka is a unified platform for handling all the real-time data feeds. Kafka supports low latency message delivery and gives guarantee for fault tolerance in the presence of machine failures. It has the ability to handle a large number of diverse consumers. Kafka is very fast, performs 2 million writes/sec. Kafka persists all data to the disk, which essentially means that all the writes go to the page cache of the OS (RAM). This makes it very efficient to transfer data from page cache to a network socket.

Kafka Cluster Architecture



Department of Computer Science and Engineering (Data Science)



Broker

Kafka cluster typically consists of multiple brokers to maintain load balance. Kafka brokers are stateless, so they use ZooKeeper for maintaining their cluster state. One Kafka broker instance can handle hundreds of thousands of reads and writes per second and each broker can handle TB of messages without performance impact. Kafka broker leader election can be done by ZooKeeper.

ZooKeeper

ZooKeeper is used for managing and coordinating Kafka broker. ZooKeeper service is mainly used to notify producer and consumer about the presence of any new broker in the Kafka system or failure of the broker in the Kafka system. As per the notification received by the Zookeeper regarding presence or failure of the broker then producer and consumer takes decision and starts coordinating their task with some other broker.

Producers

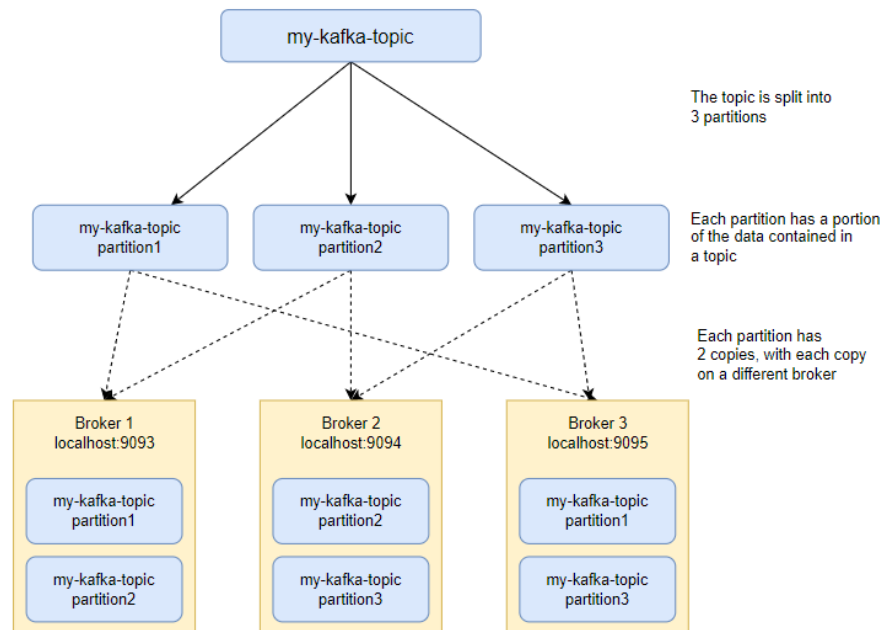
Producers push data to brokers. When the new broker is started, all the producers search it and automatically sends a message to that new broker. Kafka producer doesn't wait for acknowledgements from the broker and sends messages as fast as the broker can handle.



Department of Computer Science and Engineering (Data Science)

Consumers

Since Kafka brokers are stateless, which means that the consumer has to maintain how many messages have been consumed by using partition offset. If the consumer acknowledges a particular message offset, it implies that the consumer has consumed all prior messages. The consumer issues an asynchronous pull request to the broker to have a buffer of bytes ready to consume. The consumers can rewind or skip to any point in a partition simply by supplying an offset value. Consumer offset value is notified by ZooKeeper.



Lab Assignment:

1. Installation of Kafka 2.13-3.0.0 locally.
2. Create a Kafka local cluster with 3 brokers and create a topic to which the data should belong.
3. Send message to a topic using producer and read the data from the cluster using a consumer.
4. Testing Replication after a failed broker.



Shri Vile Parle Kelavani Mandal's

DWARKADAS J. SANGHVI COLLEGE OF ENGINEERING

(Autonomous College Affiliated to the University of Mumbai)

NAAC Accredited with "A" Grade (CGPA : 3.18)



Department of Computer Science and Engineering (Data Science)

```
cs-ds@kmaster: ~/Downloads/kafka_2.13-3.6.1
cs-ds@kmaster:~/Downloads$ tar -xzf kafka_2.13-3.6.1.tgz
cs-ds@kmaster:~/Downloads$ cd kafka_2.13-3.6.1
cs-ds@kmaster:~/Downloads/kafka_2.13-3.6.1$ bin/zookeeper-server-start.sh config/zookeeper.properties
[2024-02-21 11:03:37,156] INFO Reading configuration from: config/zookeeper.properties (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2024-02-21 11:03:37,157] WARN config/zookeeper.properties is relative. Prepend ./ to indicate that you're sure! (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2024-02-21 11:03:37,159] INFO clientPortAddress is 0.0.0.0:2181 (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2024-02-21 11:03:37,159] INFO secureClientPort is not set (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2024-02-21 11:03:37,159] INFO observerMasterPort is not set (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2024-02-21 11:03:37,159] INFO metricsProvider.className is org.apache.zookeeper.metrics.impl.DefaultMetricsProvider (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2024-02-21 11:03:37,160] INFO autopurge.snapRetainCount set to 3 (org.apache.zookeeper.server.DatadirCleanupManager)
[2024-02-21 11:03:37,160] INFO autopurge.purgeInterval set to 0 (org.apache.zookeeper.server.DatadirCleanupManager)
[2024-02-21 11:03:37,160] INFO Purge task is not scheduled. (org.apache.zookeeper.server.DatadirCleanupManager)
```



Shri Vile Parle Kelavani Mandal's

DWARKADAS J. SANGHVI COLLEGE OF ENGINEERING

(Autonomous College Affiliated to the University of Mumbai)

NAAC Accredited with "A" Grade (CGPA : 3.18)



Department of Computer Science and Engineering (Data Science)

```
cs-ds@kmaster: ~/Downloads/kafka_2.13-3.6.1
cs-ds@kmaster:~/Downloads$ bin/kafka-server-start.sh config/server.properties
bash: bin/kafka-server-start.sh: No such file or directory
cs-ds@kmaster:~/Downloads$ cd kafka_2.13-3.6.1
cs-ds@kmaster:~/Downloads/kafka_2.13-3.6.1$ bin/kafka-server-start.sh config/server.properties
[2024-02-21 11:05:32,893] INFO Registered kafka:type=kafka.Log4jController MBean (kafka.utils.Log4jControllerRegistration$)
[2024-02-21 11:05:33,065] INFO Setting -D jdk.tls.rejectClientInitiatedRenegotiation=true to disable client-initiated TLS renegotiation (org.apache.zookeeper.common.X509Util)
[2024-02-21 11:05:33,120] INFO Registered signal handlers for TERM, INT, HUP (org.apache.kafka.common.utils.LoggingSignalHandler)
[2024-02-21 11:05:33,121] INFO starting (kafka.server.KafkaServer)
[2024-02-21 11:05:33,121] INFO Connecting to zookeeper on localhost:2181 (kafka.server.KafkaServer)
[2024-02-21 11:05:33,129] INFO [ZooKeeperClient Kafka server] Initializing a new session to localhost:2181. (kafka.zookeeper.ZooKeeperClient)
[2024-02-21 11:05:33,137] INFO Client environment:zookeeper.version=3.8.3-6ad6d364c7c0bcf0de452d54ebefa3058098ab56, built on 2023-10-05 10:34 UTC (org.apache.zookeeper.ZooKeeper)
[2024-02-21 11:05:33,137] INFO Client environment:host.name=kmaster (org.apache.zookeeper.ZooKeeper)
[2024-02-21 11:05:33,137] INFO Client environment:java.version=1.8.0_392 (org.apache.zookeeper.ZooKeeper)
```



Shri Vile Parle Kelavani Mandal's

DWARKADAS J. SANGHVI COLLEGE OF ENGINEERING

(Autonomous College Affiliated to the University of Mumbai)

NAAC Accredited with "A" Grade (CGPA : 3.18)



Department of Computer Science and Engineering (Data Science)

```
cs-ds@kmaster: ~/Downloads/kafka_2.13-3.6.1
cs-ds@kmaster:~/Downloads/kafka_2.13-3.6.1$ bin/kafka-topics.sh --create --topic quickstart-events --bootstrap-server localhost:9092
Created topic quickstart-events.
cs-ds@kmaster:~/Downloads/kafka_2.13-3.6.1$ $ bin/kafka-topics.sh --describe --topic quickstart-events --bootstrap-server localhost:9092
$: command not found
cs-ds@kmaster:~/Downloads/kafka_2.13-3.6.1$ bin/kafka-topics.sh --describe --topic quickstart-events --bootstrap-server localhost:9092
Topic: quickstart-events      TopicId: dqzL_nLwRE6-GCFxgG50Dg PartitionCount: 1
ReplicationFactor: 1         Configs:
Topic: quickstart-events      Partition: 0    Leader: 0      Replicas: 0
Isr: 0
cs-ds@kmaster:~/Downloads/kafka_2.13-3.6.1$ bin/kafka-console-producer.sh --topic quickstart-events --bootstrap-server localhost:9092
>This is trinity event
>this is IDPT event
>This is hackniche
^Ccs-ds@kmaster:~/Downloads/kafka_2.13-3.6.1$ bin/kafka-console-consumer.sh --topic quickstart-events --from-beginning --bootstrap-server localhost:9092
This is trinity event
this is IDPT event
^CProcessed a total of 2 messages
cs-ds@kmaster:~/Downloads/kafka_2.13-3.6.1$ bin/kafka-console-consumer.sh --topic quickstart-events --from-beginning --bootstrap-server localhost:9092
```



Shri Vile Parle Kelavani Mandal's

DWARKADAS J. SANGHVI COLLEGE OF ENGINEERING

(Autonomous College Affiliated to the University of Mumbai)

NAAC Accredited with "A" Grade (CGPA : 3.18)



Department of Computer Science and Engineering (Data Science)

```
cs-ds@kmaster: ~/Downloads/kafka_2.13-3.6.1
cs-ds@kmaster:~/Downloads/kafka_2.13-3.6.1$ bin/kafka-console-consumer.sh --topi
c quickstart-events --from-beginning --bootstrap-server localhost:9092
This is trinity event
this is IDPT event
```




Shri Vile Parle Kelavani Mandal's

DWARKADAS J. SANGHVI COLLEGE OF ENGINEERING

(Autonomous College Affiliated to the University of Mumbai)

NAAC Accredited with "A" Grade (CGPA : 3.18)



Department of Computer Science and Engineering (Data Science)

```
cs-ds@kmaster: ~/Downloads/kafka_2.13-3.6.1
^CProcessed a total of 2 messages
cs-ds@kmaster:~/Downloads/kafka_2.13-3.6.1$ echo "plugin.path=libs/connect-file-3.6.1.jar"
plugin.path=libs/connect-file-3.6.1.jar
cs-ds@kmaster:~/Downloads/kafka_2.13-3.6.1$ echo -e "foo\nbar" > test.txt
cs-ds@kmaster:~/Downloads/kafka_2.13-3.6.1$ bin/connect-standalone.sh config/connect-standalone.properties config/connect-file-source.properties config/connect-file-sink.properties
[2024-02-21 11:27:02,964] INFO Kafka Connect worker initializing ... (org.apache.kafka.connect.cli.AbstractConnectCli:114)
[2024-02-21 11:27:02,965] INFO WorkerInfo values:
    jvm.args = -Xms256M, -Xmx2G, -XX:+UseG1GC, -XX:MaxGCPauseMillis=20, -XX:InitiatingHeapOccupancyPercent=35, -XX:+ExplicitGCInvokesConcurrent, -XX:MaxInlin
    neLevel=15, -Djava.awt.headless=true, -Dcom.sun.management.jmxremote, -Dcom.sun
    management.jmxremote.authenticate=false, -Dcom.sun.management.jmxremote.ssl=fals
    e, -Dkafka.logs.dir=/home/cs-ds/Downloads/kafka_2.13-3.6.1/bin/../logs, -Dlog4j
    configuration=file:bin/../config/connect-log4j.properties
    jvm.spec = Private Build, OpenJDK 64-Bit Server VM, 1.8.0_392, 25.392-b0
8
    jvm.classpath = /home/cs-ds/Downloads/kafka_2.13-3.6.1/bin/../libs/activ
    ation-1.1.1.jar:/home/cs-ds/Downloads/kafka_2.13-3.6.1/bin/../libs/aopalliance-r
    epackaged-2.6.1.jar:/home/cs-ds/Downloads/kafka_2.13-3.6.1/bin/../libs/argparse4
    j-0.7.0.jar:/home/cs-ds/Downloads/kafka_2.13-3.6.1/bin/../libs/audience-annotati
    ons-0.12.0.jar:/home/cs-ds/Downloads/kafka_2.13-3.6.1/bin/../libs/caffeine-2.9.3
```




Shri Vile Parle Kelavani Mandal's

DWARKADAS J. SANGHVI COLLEGE OF ENGINEERING

(Autonomous College Affiliated to the University of Mumbai)

NAAC Accredited with "A" Grade (CGPA : 3.18)



Department of Computer Science and Engineering (Data Science)

```
cs-ds@kmaster: ~/Downloads/kafka_2.13-3.6.1
[2024-02-21 11:28:16,345] WARN [Consumer clientId=console-consumer, groupId=console-consumer-7929] Error while fetching metadata with correlation id 2 : {connect-test=LEADER_NOT_AVAILABLE} (org.apache.kafka.clients.NetworkClient)
echo Another line>> test.txt
^CProcessed a total of 0 messages
schema:type:string: command not found
schema:type:string: command not found
...: command not found
cs-ds@kmaster:~/Downloads/kafka_2.13-3.6.1$ echo -e "foo\nbar"
foo
bar
cs-ds@kmaster:~/Downloads/kafka_2.13-3.6.1$ echo -e "foo\nbar" test.txt
foo
bar test.txt
cs-ds@kmaster:~/Downloads/kafka_2.13-3.6.1$ more test.txt
foo
bar
cs-ds@kmaster:~/Downloads/kafka_2.13-3.6.1$ bin/kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic connect-test --from-beginning
{"schema":{"type":"string","optional":false},"payload":"foo"}
{"schema":{"type":"string","optional":false},"payload":"bar"}
...
echo Another line>> test.txt
```

```
connect-standalone.properties  test.txt
: foo
: bar
> Another line
```